

# TEXT AS DATA: WEEK 9

MATTHIAS HABER

10 NOVEMBER 2021

# GOALS FOR TODAY

# GOALS

- Scaling Models with Wordscores
- Scaling Models with Wordfish

# WORDSCORES

# WORDSCORES

Wordscores compares the word frequencies of texts at hand to the word frequencies of so called reference texts with known (or assumed) positions and assigns document scores based on the similarity of these references.

- Highly automated, (nearly) no language knowledge needed

# WORDSCORES CONCEPT

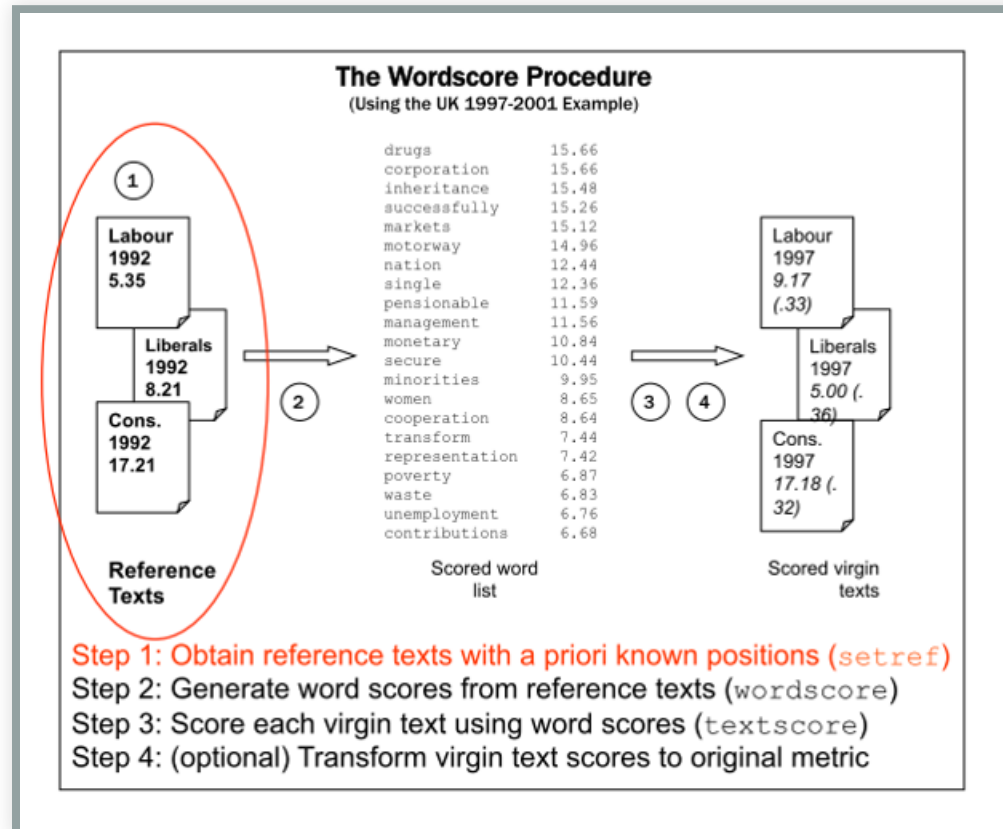
## The idea

- Each word has a policy position (word score)
- Some reference document positions are known
- Document positions are average of its words' positions
- 1st step: Derive wordscores from reference texts
- 2nd step: Apply wordscores to virgin texts

# WORDSCORES: ILLUSTRATIVE EXAMPLE

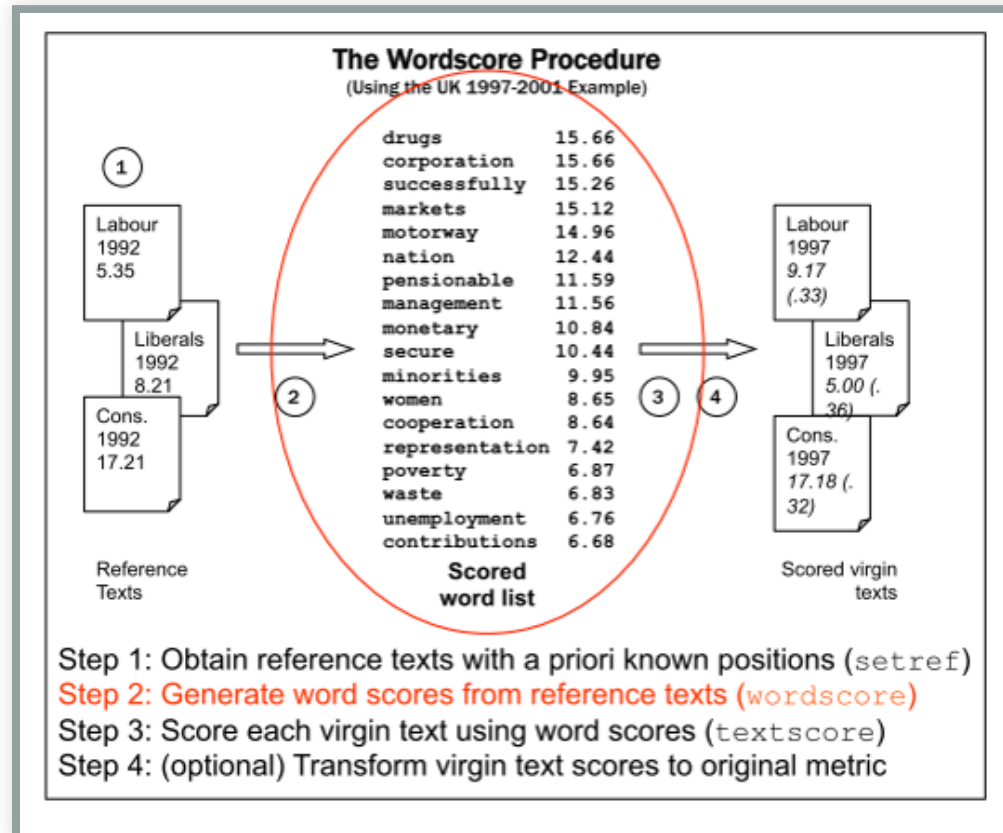
- Consider two reference texts A and B
- The word “choice” is used 10 times per 10,000 words in text A and 30 times per 10,000 words in text B
- Conditional on observing the “choice”, we are reading text A with probability 0.25 and text B with 0.75
- We can compute a “word score” once we assign reference values to the reference texts
- Suppose reference text A has position -1, and text B position +1
- then the score of word “choice” is:
- $0.25(-1.0) + 0.75(1.0) = -0.25 + 0.75 = 0.5$

# WORDSCORES PROCEDURE

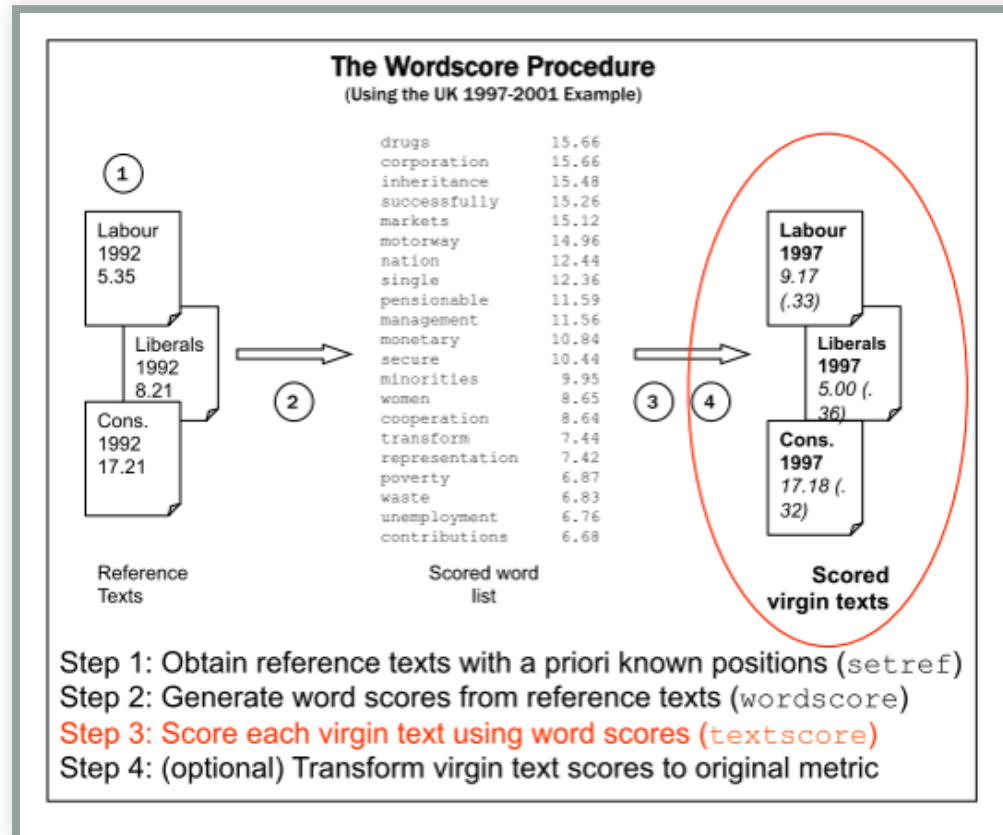




# WORDSCORES PROCEDURE



# WORDSCORES PROCEDURE



# ESTIMATING POLICY POSITIONS FROM POLITICAL TEXTS

## Laver & Garry

- Goal: Generating party positions for British and Irish manifestos
- Coding scheme similar to the CMP's
  - More hierarchical, larger number of categories
  - Each category has a pro-, con- and neutral variant

# ESTIMATING POLICY POSITIONS FROM POLITICAL TEXTS

## Assumptions:

1. Manifesto content is related to party policy positions
2. Word usage is related to policy positions
3. Word usage is constant over time
4. All relevant words are covered in the reference texts

# ESTIMATING POLICY POSITIONS FROM POLITICAL TEXTS

## 1st step: Training set

- Manifestos of Labour and Cons (UK) in 1992
  - Pool of 'keywords'
  - $N_L \geq 2N_R \Rightarrow$  Dictionary element left
  - $N_R \geq 2N_L \Rightarrow$  Dictionary element right
- Allocate selected words to the coding scheme's categories

# ESTIMATING POLICY POSITIONS FROM POLITICAL TEXTS

2nd step: Count occurrence of elements in the dictionary in manifestos

- Britain (1992 & 1997)
- Ireland (1992 & 1997)
- Left-right-scaling:  $\frac{R-L}{R+L}$ 
  - $Econ_{LR}$
  - $Soc_{LR}$

# ESTIMATING POLICY POSITIONS FROM POLITICAL TEXTS

## Test-Set: Crossvalidation

- Expert Surveys
- CMP Coding/Revised CMP Coding

**TABLE 3** Pearson Correlations between Alternative Estimates of Economic Left-Right Scale Positions, Britain and Ireland 1992–97

	Computer Codings	Revised Expert Codings	Original MRG Codings	Expert Surveys
1992				
Computer codings	1.00			
Revised expert codings	0.85	1.00		
Original MRG codings	0.72	0.94	1.00	
Expert surveys	0.75	0.95	0.99	1.00
1997				
Computer codings	1.00			
Revised expert codings	0.94	1.00		
Expert surveys	0.91	0.95	n.a	1.00

# WORDSCORES AND DICTIONARIES

- Conceptually, the two steps do the same in both approaches:
  - 1st step derives a position of a word from texts with known properties
  - 2nd step weighs the words in the unknown texts with this information
- Information in dictionary is often binary, wordscores in wordscore are scale



# SELECTING REFERENCE TEXTS

- Reference texts should use the same vocabulary in the same context
- Reference texts need to span the full dimension
- Set of reference text should contain as many words as possible
- Estimates of the positions (reference scores) need to be well grounded and/or very conservative

# 1ST STEP - OBTAINING WORDSCORES

- Start out from the observed word frequencies in reference texts:
- $F_{wr}$ : Relative frequency of word  $w$  in reference-text  $r$
- Conditional probabilities: Given we are observing word  $w$ , what is the probability that we are reading text  $r$ ?
- $$P_{wr} = \frac{F_{wr}}{\sum_r F_{wr}}$$
- $$S_w = \sum_r (P_{wr} * A_r)$$
  - $A_r$  is the a priori score for reference text  $r$
  - $S_w$  is the actual wordscore for word  $w$

## 2ND STEP – APPLYING THE WORDSCORES

- $S_v = \sum_w (F_{wv} * S_w)$ 
  - $F_{wv}$  is analogous to  $F_{wr}$
  - $S_v$  is the weighted mean score of the words in text  $v$
- Variance is the basis for calculating uncertainty
  - Summary for the consensus of the scores of each word in the virgin text
  - Higher consensus -> lower variance -> less uncertainty

# WORDSCORES – PRACTICAL CONSIDERATIONS

- What about scaling things other than manifestos?
- E.g. speeches:
  - Use reference texts from other context (e.g. manifesto)?
  - What scores to use?
  - Length of the reference texts?


## COMPARING WORDSCORES AND CMP

- Laver, Garry & Benoit use Irish and British manifestos to demonstrate/validate
- CMP data offers information for many countries and over long periods
- How do wordscores results compare across countries?


# COMPARING WORDSCORES AND CMP

- Bräuninger, Debus & Müller (2013) compare wordscores results for 13 countries between 1980 and 2000
- Reference texts are the manifestos in the latest elections
- Reference scores are the Rile scores from CMP
- Essentially 'replicate' CMP scores using wordscores

# COMPARING WORDSCORES AND CMP



# COMPARING WORDSCORES AND CMP





# COMPARING WORDSCORES AND CMP

- Wordscores replicates CMP better when
  - reference texts cover the full range of a dimension
  - the percentage of scored words is high
- Cross-check results from wordscores before using them in an analysis

# WORDSCORES EXERCISE: LOAD DATA AND CREATE CORPUS

In the wordscores exercise we'll be looking at annual budget speeches held in the Irish Parliament from 2008 - 2012. Let's load the data, tokenize it and create a dfm.

```
library(quanteda.textmodels)
data(data_corpus_irishbudget2010, package = "quanteda.textmodels")
budget_dfm <- dfm(tokens(data_corpus_irishbudget2010))
```

# WORDSCORES EXERCISE: LOAD DATA AND CREATE CORPUS

As we just learned wordscores requires us to assign a set of known scores to so called reference texts to identify the positions of new documents. Let's create some reference scores for the 5th and 6th document.

```
refscores <- c(rep(NA, 4), 1, -1, rep(NA, 8))
```

# WORDSCORES EXERCISE: PREDICT WORD POSITIONS

We can use the `textmodel_wordscores ( )` function from the `quanteda.textmodels` package to estimate the document positions based on the positions of the reference texts.

```
ws <- textmodel_wordscores(budget_dfm, y = refscores, smooth = 1)
head(coef(ws))
```


```
##           when           i      presented           the supplementary
## -0.53014011  -0.28805749  -0.42939762    0.01670114  -0.42939762
##           budget
## -0.15483456
```

# WORDSCORES EXERCISE: PLOT WORD POSITIONS

We can use the `textplot_scale1d()` function from the `quanteda.textplots` package to plot word positions.

```
library(quanteda.textplots)
textplot_scale1d(ws,
                 highlighted = c("cut", "minister", "welfare",
                                "economy", "budget"),
                 highlighted_color = "red")
```

# WORDSCORES EXERCISE: PLOT WORD POSITIONS



# WORDSCORES EXERCISE: PREDICT DOCUMENT POSITIONS

We can extract the estimated speaker positions using the `predict ( )` function as if this were a regular fitted regression model.

```
ws_pred <- predict(ws, interval = "confidence")  
ws_pred
```

```
## $fit  
##  
##           fit           lwr           upr  
## Lenihan, Brian (FF)      0.116362873  0.110173633  0.122552114  
## Bruton, Richard (FG)     0.003318195 -0.005228996  0.011865387  
## Burton, Joan (LAB)      -0.004696791 -0.011996402  0.002602821  
## Morgan, Arthur (SF)     0.019334198  0.012308623  0.026359774  
## Cowen, Brian (FF)       0.213090960  0.206424497  0.219757423  
## Kenny, Enda (FG)        -0.136776116 -0.146311207 -0.127241025  
## ODonnell, Kieran (FG)   -0.008132150 -0.020643605  0.004379306  
## Gilmore, Eamon (LAB)    0.013988095  0.004862201  0.023113990  
## Higgins, Michael (LAB)  -0.019145477 -0.035788594 -0.002502360  
## Quinn, Ruairi (LAB)     -0.027387712 -0.043704917 -0.011070506  
## Gormley, John (Green)    0.085698169  0.067187907  0.104208431  
## Ryan, Eamon (Green)     0.056610204  0.042201917  0.071018490  
## Cuffe, Ciaran (Green)   0.072391730  0.055483988  0.089299472  
## OCaolain, Caoimhghin (SF) 0.016769698  0.007557775  0.025981621
```


# WORDSCORES EXERCISE: PLOT SPEAKER POSITIONS

And then plot the results using the same `textplot_scale1d()` function and grouping by party.

```
textplot_scale1d(ws_pred, margin = "documents",  
                 groups = docvars(data_corpus_irishbudget2010,  
                                  "party"))
```



# WORDSCORES EXERCISE: PLOT SPEAKER POSITIONS



# WORDFISH

# WORDSCORES VS. WORDFISH

- Wordscores derives and 'transfers' known properties of words, i.e. the wordscores between texts.
- Wordfish builds a statistical model that explains the occurrence of each word: Poisson regression

# WORDSCORES VS. WORDFISH

- Advantages from a practical perspective
  - No reference texts needed; Anchor points instead
  - Statistically models all words in a text
  - Absolute minimum of input from the user; Versatile and well suited for smaller projects
- The statistical model replaces the need for reference texts
  - Mathematical complexity of the model

# POISSON MODEL

- Dependent variables of interest may be counts, e.g.
  - Occurrence of conflict/wars, casualties in conflicts; Number of bills brought forward in a term; Number of hospitalizations, sicknesses etc.
  - **Word count in a document**
- A dependent count variable  $\gamma$ 
  - bound between 0 and  $\infty$
  - takes only discrete values (0,1,2,3,...)

# POISSON MODEL

## Poisson distribution

$$y_i = \text{Poisson}(\lambda_i)$$

- Poisson distribution: Repeated Bernoulli-Experiments (0/1)
- Generally used in count data (poisson regression)
- Has only one parameter: 'Event occurrence rate'
- No contagion effects; the event rate remains constant

# WORDFISH MODEL

$$\gamma_{ij} \sim \text{Poisson}(\lambda_{ij})$$
$$\lambda_{ij} = \exp(\alpha_i + \psi_j + \beta_j * \omega_i)$$

- $i$  = document (e.g. party manifesto)
- $j$  = unique word
- $\alpha_i$  = document fixed effect
- $\psi_j$  = word fixed effect
- $\beta_j$  = word specific weight (sign representing the ideological direction)
- $\omega_i$  = document position

# WORDFISH ESTIMATION

- Regression without independent variables
    - Solution: Maximum Likelihood Estimation
1. Estimate party parameters conditional on the expectation for the word parameters (in first iteration the starting values)
  2. Estimate word parameters conditional on party parameters obtained in previous step
  3. Go back and forth until a convergence criterion is met and the likelihoods do not change anymore



# WORDFISH ESTIMATION

## Likelihood function

$$\sum_j^m \sum_i^n -\exp(\alpha_i + \psi_j + \beta_j * \omega_i) + \ln(\exp(\alpha_{it} + \psi_j + \beta_j * \omega_i))$$

Without fixing some parameters, there are infinite combinations of  $\omega$  and  $\beta$ , which could provide the same likelihood.

- Document positions: mean of all positions  $\omega$  across all elections is set to 0, and standard deviation to 1.
- Set directionality (e.g. document A always has a smaller valuer than document B).
- Set document fixed effect: first document  $\alpha$  is set to 0.

## WORDFISH ESTIMATION

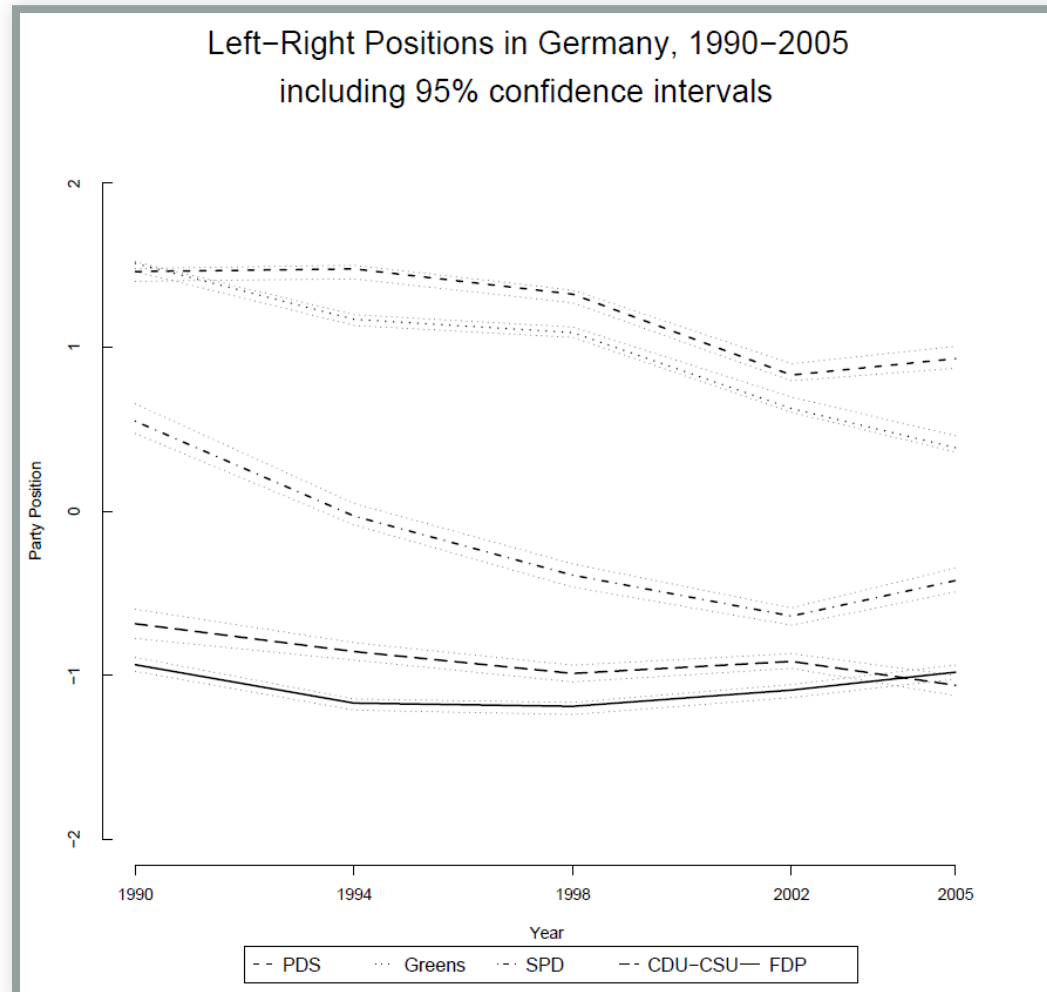
This means that you **cannot directly compare estimates ACROSS different estimations** (for example, in secondary analysis). This is the case for ALL scaling models (e.g. Nominate), and also for Wordscores.

- Think to what extent position estimates are actually comparable...
  - ... across countries
  - ... over time
  - ... between documents


# WORDFISH ESTIMATION

- Dimension of the scaling is created ex post (as compared to Wordscores)
  - What is the dimension identified?
  - More validation required
  - Creation of alternative dimensions via subsetting texts only

# WORDFISH OUTPUT



# WORDFISH OUTPUT




# WORDFISH: MULTILINGUAL?

Does it work in different languages? What are some reasons for doubting it?

- Ideal case: get the exact same political texts in high quality translations
- Estimate Wordfish and compare across different languages
- This is possible: European Parliament speeches (translated into all official languages)

# WORDFISH: SPEECHES IN EU PARLIAMENT

Positions of National Parties in the European Parliament



Source: Proksch and Slapin (2010)

# (IN)STABILITY OF THE POLITICAL LEXICON

- What if the political lexicon is unstable over time?
  - New issues appear, old issues disappear
- If this happens frequently, then scaling algorithms will pick up shifts in the policy agenda – rather than shifts in party positions.
- In fact, this is one assumption: that word usage reflects ideology.
  - For example, it becomes seriously problematic when all parties start talking about the “issue” of the day. Then we can distinguish between elections, but not very well between parties



# POLICY DIMENSIONALITY

- In the original Wordscores article, the authors assumed the policy dimensions can be chosen by using different reference scores from known policy scales (expert surveys)
  - “Economic policy” and a “social policy” for the UK in 1997
  - Reference texts: 3 manifestos from 1992 (same for BOTH dimensions)

Reference Scores	Lib Dems	Labour	Cons.
Economic Policy	8.21	5.35	17.21
Social Policy	6.87	6.53	15.34

# POLICY DIMENSIONALITY

- This means that estimates across dimensions vary only because different expert evaluations are used to anchor the texts, not because the text input varies according to the policy area of interest
  - This is not a necessary assumption. There is no reason why one cannot run analysis of sort of a policy dictionary.
  - One possibility is to preserve as much of the context as possible and parse the sections of manifestos into policy areas and then estimate positions on those sections only (Proksch and Slapin 2006, Slapin and Proksch 2008)

# WORDFISH EXERCISE: LOAD DATA

We will take a look at the US Senate debate on partial birth abortion.

First let's load the data into R and take a look at the corpus:

```
library(quanteda)
load("data/corpus_us_debate_speaker.rda")
summary(corpus_us_debate_speaker)
```

```
## Corpus consisting of 23 documents, showing 23 documents:
```

```
##
```

```
##      Text Types Tokens Sentences party      speaker
```

```
##      ALLARD   400   1165         53      R      ALLARD
```

```
##      BOND    129    232          9      R      BOND
```

```
##      BOXER   2231  18527        886      D      BOXER
```

```
##      BROWNB  646   2884        168      R      BROWNB
```

```
##      BUNNING 281    593         32      R      BUNNING
```

```
##      CANTWELL 395   1114         55      D      CANTWELL
```

```
##      DeWINE  463   1438         75      R      DeWINE
```

```
##      DOMENICI 203    479         27      R      DOMENICI
```

```
##      DURBIN  520   1874         89      D      DURBIN
```

```
##      ENSIGN  410   1235         66      R      ENSIGN
```

```
##      FEINGOLD 213    414         19      D      FEINGOLD
```

```
##      FEINSTEIN 1248  6112        284      D      FEINSTEIN
```

##	FEINSTEIN	1348	8112	284	D	FEINSTEIN
##	FRIST	222	501	25	R	FRIST
##	HARKIN	655	2612	145	D	HARKIN
##	HATCH	451	1173	60	D	HATCH

# WORDFISH EXERCISE: TOKENIZE AND DFM

Now we'll tokenize the text and create a document term matrix.

```
stops <- c(setdiff(stopwords(),  
                  c("her", "she", "hers", "he", "his", "him")),  
          "bill", "can")  
  
toks <- tokens(corpus_us_debate_speaker,  
              remove_punct = TRUE,  
              remove_symbols = TRUE,  
              remove_numbers = TRUE) |>  
  tokens_remove(stops) |>  
  tokens_tolower()  
  
senate_dfm <- dfm(toks)
```

## WORDFISH EXERCISE: RUN THE WORDFISH MODEL

Let's run the one-dimensional scaling model known as wordfish using the `textmodel_wordfish` function from the `quanteda.textmodels` package.

```
library(quanteda.textmodels)
wf <- textmodel_wordfish(senate_dfm, dir = c(3, 21))
```

We use the `dir` argument to set the polar opposites of the scale. In this case we assume that the 3rd speaker Barbara Boxer, is to the left of the 21st speaker Rick Santorum.

# WORDFISH EXERCISE: INVESTIGATE THE OUTPUT

Let's take a look at the components of the fitted model.

```
wf_coef <- coef(wf)
names(wf_coef)
```

```
## [1] "documents" "features"
```

```
head(wf_coef$documents)
```

```
##           theta      alpha
## ALLARD      0.02165274  0.21348639
## BOND       -0.08693357 -1.29227632
## BOXER      -1.05777079  2.76206193
## BROWNBAC  2.38269597  0.12975044
## BUNNING     0.81776695 -0.48868033
## CANTWELL   -1.25634814 -0.07911048
```

```
head(wf_coef$features)
```

```
##           beta      psi
## mr          0.06468449  0.8700898
## "            0.05468101  1.5000171
```

```
## president      0.05462191  1.5038171
## commend        -0.02595502 -2.0448676
## senator         0.06740927  1.8584309
## pennsylvania   -0.13413832 -0.6192956
## santorum        0.12538330 -0.4688482
```



# WORDFISH EXERCISE: INVESTIGATE THE OUTPUT

```
head(wf_coef$features)
```

```
##              beta      psi
## mr           0.06468449  0.8700898
## president    0.05462191  1.5038171
## commend      -0.02595502 -2.0448676
## senator       0.06740927  1.8584309
## pennsylvania -0.13413832 -0.6192956
## santorum      0.12538330 -0.4688482
```

# WORDFISH EXERCISE: ESTIMATE SPEAKER POSITIONS

We can extract the estimated speaker positions using the `predict ( )` function just as with wordscores.

```
preds <- predict(wf, interval = "confidence")
preds
```

```
## $fit
##           fit           lwr           upr
## ALLARD      0.02165274 -0.10190180  0.1452073
## BOND       -0.08693357 -0.34076934  0.1669022
## BOXER      -1.05777079 -1.07929641 -1.0362452
## BROWNBAC  2.38269597  2.28117382  2.4842181
## BUNNING     0.81776695  0.62356627  1.0119676
## CANTWELL   -1.25634814 -1.33174776 -1.1809485
## DeWINE     1.38728909  1.25204737  1.5225308
## DOMENICI    0.32865152  0.12357508  0.5337280
## DURBIN     -0.25969664 -0.35154658 -0.1678467
## ENSIGN     0.83658945  0.68738863  0.9857903
## FEINGOLD   -1.20437761 -1.32919832 -1.0795569
## FEINSTEIN  -1.59354288 -1.61938317 -1.5677026
## FRIST      0.83880779  0.61988997  1.0577256
## HARKIN     -0.79094571 -0.85650252 -0.7253889
## HATCH      0.97739715  0.83244576  1.1223485
## LAUTENBERG -0.65712119 -0.73253786 -0.5817045
```

# WORDFISH EXERCISE: CREATE A NICE DATA FRAME

If we want to we can extract the estimated speaker positions into a data frame and merge in the docvars.

```
speaker_pos <- data.frame(preds$fit,  
                           speaker = row.names(preds$fit)) %>%  
  dplyr::left_join(docvars(senate_dfm)) %>%  
  dplyr::arrange(fit)  
head(speaker_pos)
```


##		fit	lwr	upr	speaker	party
##	1	-1.593543	-1.619383	-1.5677026	FEINSTEIN	D
##	2	-1.256348	-1.331748	-1.1809485	CANTWELL	D
##	3	-1.226258	-1.325300	-1.1272152	MIKULSKI	D
##	4	-1.204378	-1.329198	-1.0795569	FEINGOLD	D
##	5	-1.102293	-1.258032	-0.9465541	NELSO	D
##	6	-1.057771	-1.079296	-1.0362452	BOXER	D

# WORDFISH EXERCISE: LET'S PLOT THE SPEAKER POSITIONS

We can use the `textplot_scale1d()` function from the `quanteda.textplots` package to plot speaker positions.

```
library(quanteda.textplots)
textplot_scale1d(wf, groups = corpus_us_debate_speaker$party)
```

# WORDFISH EXERCISE: LET'S PLOT THE SPEAKER POSITIONS



# WORDFISH EXERCISE: LET'S PLOT THE WORD POSITIONS

Similarly, we can use the same function to plot the estimated positions of each word.

```
textplot_scaleld(wf, margin = "features", alpha = 0.2,  
                 highlighted = c("life", "choice", "womb", "her",  
                                "woman", "health",  
                                "born", "baby", "gruesome", "kill",  
                                "roe", "wade",  
                                "medical", "her", "his", "child",  
                                "religion", "doctor"),  
                 highlighted_color = "red")
```

The word cloud plot displays the relationship between estimated beta (x-axis) and estimated psi (y-axis) for various terms. The x-axis ranges from -5 to 5, and the y-axis ranges from -7.5 to 2.5. Words are colored based on their estimated psi value, with a color scale from blue (low psi) to red (high psi). The plot shows a dense cluster of words in the center, with some words appearing more frequently than others. The words are arranged in a way that suggests a positive correlation between estimated beta and estimated psi, with higher psi values generally corresponding to higher beta values.

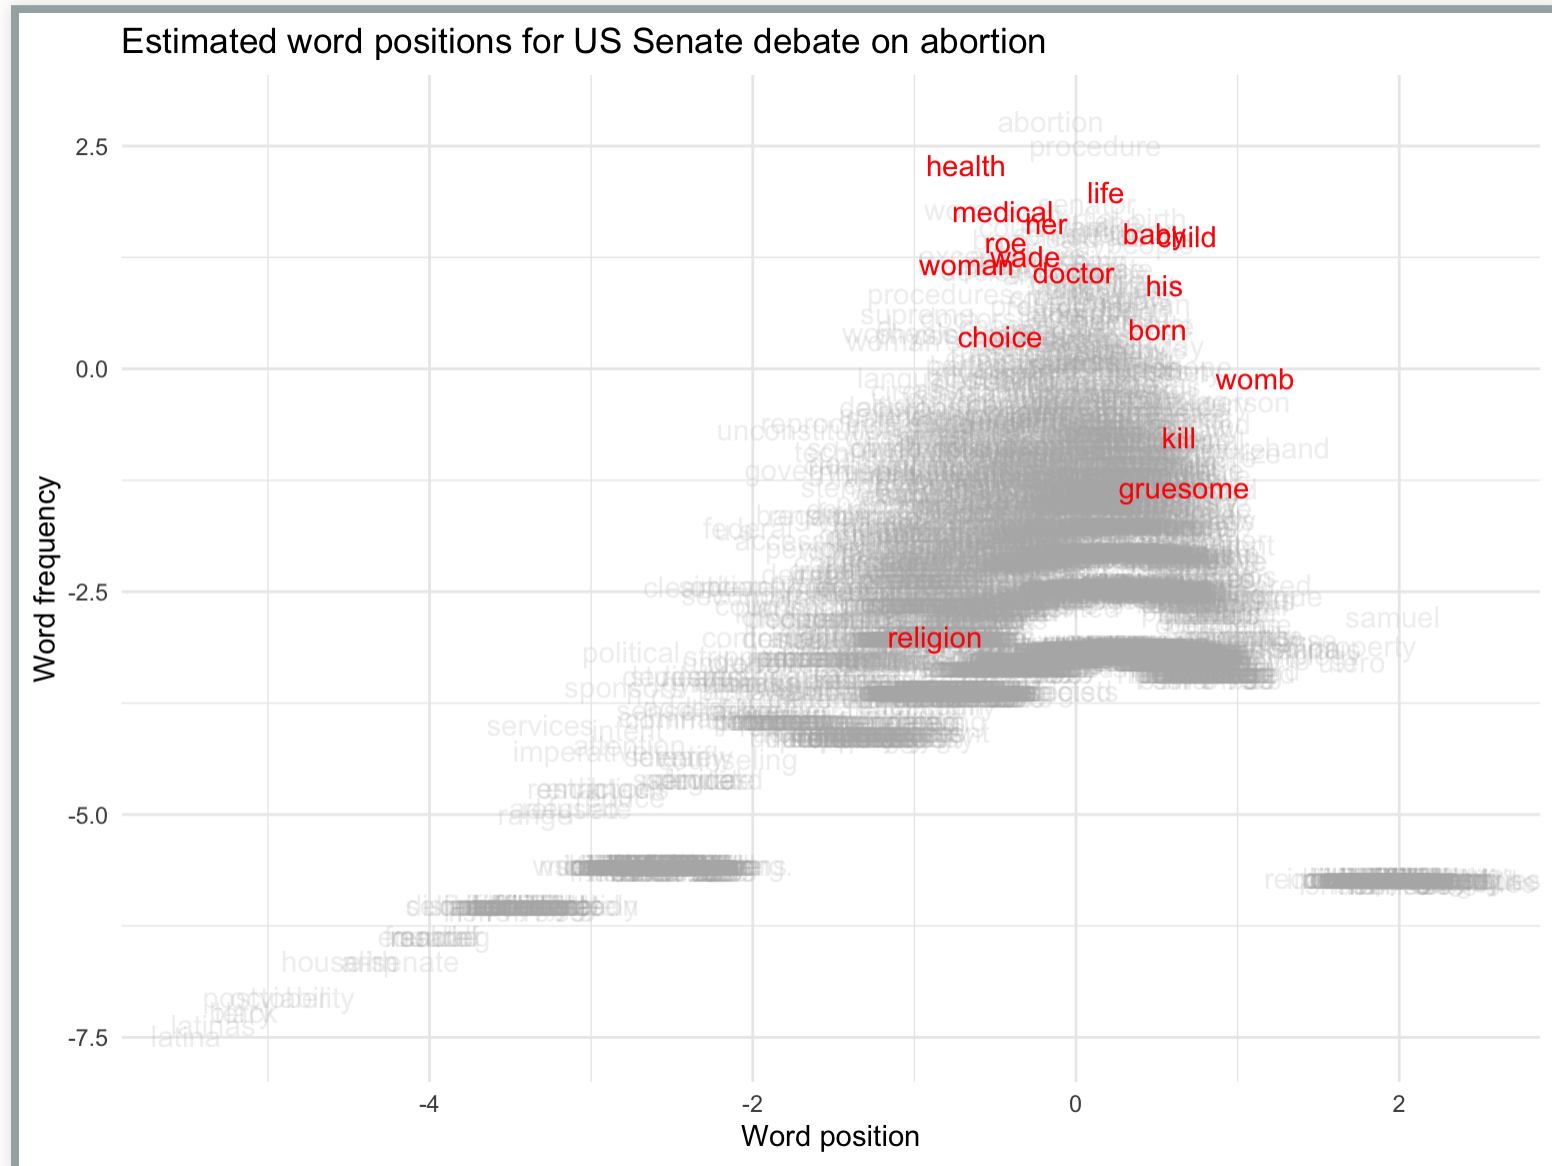
# WORDFISH EXERCISE: LET'S PLOT THE WORD POSITIONS

We can also use our ggplot syntax to make the plot a bit nicer.

```
textplot_scaled(wf, margin = "features", alpha = 0.2,  
                highlighted = c("life", "choice", "womb", "her",  
                                "woman", "health",  
                                "born", "baby", "gruesome", "kill",  
                                "roe", "wade",  
                                "medical", "her", "his", "child",  
                                "religion", "doctor"),  
                highlighted_color = "red") +  
labs(x = "Word position", y = "Word frequency") +  
ggtitle("Estimated word positions for use senate debate on abortion")  
+  
theme_minimal()
```



# WORDFISH EXERCISE: LET'S PLOT THE WORD POSITIONS



# HOMEWORK

Create a corpus from a number of documents of your own choosing and use either the wordscores or the wordfish approach to estimate the document positions. Plot your results and share them in class next week.

# WRAPPING UP

# QUESTIONS?

# OUTLOOK FOR OUR NEXT SESSION

Next week we will look at different ways of generating topics from text.

# THAT'S IT FOR TODAY

Thanks for your attention!

